

Single-Deploy Modular Monorepo (SDMM) Engineering

Principles & Guidelines

AzzCraft, Inc. (艾之舟科技)

January 2026

目录

Normative language	4
1 Purpose	4
2 Non-goals	4
3 Definitions	5
3.1 App shell and composition root	5
3.2 Package	5
3.3 Contract	5
3.4 Boundary	6
4 Core reasoning: efficiency - when SDMM helps vs hurts	6
4.1 Why SDMM usually avoids “microservice-style” efficiency loss	6
4.2 What can regress runtime performance in SDMM	6
4.3 What can regress developer/CI efficiency in SDMM	6
5 Reference architecture	7
5.1 Workspace layout	7
5.2 Package roles - responsibility boundaries	7
6 Foundational principles	8
6.1 P1 — Contracts first, implementation second	8
6.2 P2 — Single owner per cross-cutting concern	8
6.3 P3 — Features own user flows, not infrastructure	9
6.4 P4 — Acyclic dependency graph, enforced	9

⁰© 2026 AzzCraft, Inc. (艾之舟科技) All rights reserved. This document is published as a public reference. Commercial use, redistribution, or derivative works require prior written permission. For licensing inquiries, contact qipeng@azzcraft.com.

6.5	P5 — Public API only - no deep imports	11
6.6	P6 — Backward-compatible behavior is the default	11
6.7	P7 — Make packages small, cohesive, and replaceable	12
6.8	P8 — The app shell is the composition root	12
6.9	P9 — Prefer explicit dependency injection over hidden globals	12
7	Contract rules	13
7.1	Contract sources of truth	13
7.2	Contract evolution: refactors vs versioning	13
7.3	Contract drift prevention	13
8	api package guidelines (I/O layer)	14
8.1	Responsibilities	14
8.2	Configuration injection	14
8.3	Timeouts, cancellation, and long-running operations	16
8.4	Error model ownership	16
9	contracts package guidelines (contract layer)	17
9.1	What belongs here	17
9.2	What must not belong here	18
10	core package guidelines (shared policy + pure logic layer)	18
10.1	Responsibilities	18
10.2	Keep <code>core</code> deterministic by default	18
10.3	Framework specificity	19
11	ui package guidelines (shared presentation layer)	19
11.1	Responsibilities	19
11.2	Rules (avoid hidden coupling)	19
11.3	Dependency rule: do not couple <code>ui</code> to <code>core</code>	19
11.4	Styling and assets strategy (required)	20
12	Feature package guidelines (features/*)	20
12.1	Feature boundaries	20
12.2	Feature public exports	21
13	App shell guidelines (apps/*)	21
14	State management principles	21
14.1	Where state lives	21
14.2	Avoid hidden coupling	21

15 Performance and build efficiency principles	22
15.1 Guarantee singleton framework/runtime dependencies	22
15.2 Deliberate code splitting	22
15.3 Export hygiene (avoid barrel bloat)	22
15.4 Build/CI efficiency	22
16 Error handling, security, and observability	22
16.1 One global error policy	22
16.2 Correlation IDs and telemetry	22
16.3 Secret handling	23
16.4 Untrusted inputs and boundary validation	23
17 Testing strategy	23
17.1 What to test where	23
17.2 Contract tests	23
18 Tooling enforcement	23
19 AI-assisted development workflow	24
20 Decision guide	24
20.1 S1 — Add a new screen/page	24
20.2 S2 — Add a new backend/API endpoint	24
20.3 S3 — Add a shared UI component	24
20.4 S4 — Add shared logic/utility	24
20.5 S5 — Add a new permission or gated UI action	24
20.6 S6 — Add a streaming/realtime feature	25
20.7 S7 — Add a long-running compute/generation operation	25
20.8 S8 — Two features need the same domain type	25
20.9 S9 — Feature A needs data produced by Feature B	25
20.10S10 — Change a contract (shape/semantics/identifier)	25
20.11S11 — A feature package becomes too large	25
20.12S12 — You need to break a dependency rule	25
Appendix A. Migration playbook	26
21 SDMM checklist	26

Status: General-purpose reference for implementation and AI-assisted development
Scope: Any codebase using workspace packages (pnpm/yarn/npm workspaces, Nx/Turbo optional) to split a frontend into small, contract-driven packages while keeping each deployed frontend as a **single runtime artifact**.

Normative language

The key words **MUST**, **MUST NOT**, **SHOULD**, **SHOULD NOT**, and **MAY** are to be interpreted as described in RFC 2119.

1 Purpose

SDMM is a build-time modularization pattern:

- Split a codebase into **small, independently editable packages** (API, contracts, shared logic, UI, features).
- Keep each product frontend as **one deployed application** (one SPA/SSR artifact), avoiding micro-frontend runtime complexity.
- Make boundaries explicit so that **AI agents** can implement isolated changes safely without breaking integration.

SDMM is not a “folder refactor”. It is a **boundary-and-contract refactor** that aims to reduce:

- accidental coupling
 - inconsistent error/auth behavior
 - bundle bloat from unmanaged shared code
 - integration regressions caused by unclear ownership
-

2 Non-goals

SDMM is **not**:

- 1) A micro-frontend architecture (separate deployments + runtime composition).

⁰© 2026 AzzCraft, Inc (艾之舟科技) . All rights reserved. This document is published as a public reference. Commercial use, redistribution, or derivative works require prior written permission. For licensing inquiries, contact qipeng@azzcraft.com.

- 2) A justification to duplicate domain logic in multiple packages.
 - 3) A replacement for contract tests and end-to-end validation.
 - 4) A guarantee of faster builds: SDMM can slow builds if boundaries are unstable or tooling is weak.
-

3 Definitions

3.1 App shell and composition root

- **App shell:** the runnable application entry for a deployable artifact (e.g., `apps/web`). It owns bootstrap, routing, global providers, and cross-package composition.
- **Composition root:** the place where the app chooses implementations and wires dependencies.

In SDMM there is **one app shell per deployed artifact**. A single repo MAY contain multiple app shells (e.g., `apps/customer-web`, `apps/admin-web`). Each app shell is still “single-deploy”.

3.2 Package

A **package** is a workspace module under `packages/*` that is:

- build-time linked into an app shell build
- versioned and reviewed as an independent unit of change
- small enough that an AI agent can safely modify it with minimal context

Packages do not imply separate runtime deployments.

3.3 Contract

A **contract** is anything multiple packages (or systems) must agree on. Contracts include:

- data: request/response DTO shapes, error payload shapes, event/message schemas
- identifiers: route IDs, permission codes, feature flags
- semantics: auth/session expectations, retry/timeout policy, pagination behavior, streaming lifecycle rules

Important distinction:

- A **contract (concept)** can include semantics and rules.
- A **contracts package (code)** MUST contain only the **representations** of contracts (types/schemas/constants), not the behavioral implementations.

Behavior that enforces a contract lives in `api / core` / app shell and is validated via tests.

3.4 Boundary

A **boundary** is the enforced line between packages. Boundaries are defined by:

- allowed imports
- public exports
- ownership and responsibility

SDMM only works when boundaries are real and enforced.

4 Core reasoning: efficiency - when SDMM helps vs hurts

4.1 Why SDMM usually avoids “microservice-style” efficiency loss

SDMM typically does **not** add runtime latency because it does **not** introduce:

- extra network hops
- extra backend calls
- separate runtime deployments of frontend parts

Most costs are build-time (tooling, packaging) and organizational (ownership discipline).

4.2 What can regress runtime performance in SDMM

Runtime regressions occur when the split causes:

- **duplicate runtime copies** (two Reacts, two Vue runtimes, duplicate UI library copies)
- loss of tree-shaking due to careless exports (“barrel bloat”)
- accidental de-optimization of code splitting (everything becomes eagerly imported)
- duplicated large assets/fonts or repeated CSS injection

4.3 What can regress developer/CI efficiency in SDMM

Developer/CI regressions occur when:

- boundaries are unstable (constant cross-package churn)
 - “shared” packages become dumping grounds (**core** grows without discipline)
 - import boundaries are not enforced (entropy: “monolith with folders”)
 - tasks are not cached (no incremental build/test strategy)
 - packages have inconsistent tooling versions/configs
-

5 Reference architecture

5.1 Workspace layout

```
repo-root/
  apps/
    web/                      # app shell (single deploy)
    # other app shells may exist; each is single-deploy

  packages/
    contracts/                # shared types + schemas + contract constants
    foundation/               # OPTIONAL: shared pure runtime helpers (below core/ui)
    api/                      # network I/O: clients + endpoint wrappers
    core/                     # shared non-visual logic + policies (I/O-free by default)
    ui/                       # reusable UI components + design primitives

    features/
      feature-a/
      feature-b/
      ...

```

`foundation/` is optional but strongly recommended if you want **shared helpers** without creating a `ui -> core` dependency.

5.2 Package roles - responsibility boundaries

- **contracts**
 - Owns shared **types/schemas/constants** for contracts.
 - MUST remain dependency-light.
- **foundation (optional)**
 - Owns shared **pure runtime utilities** used across layers:
 - * formatting helpers that are truly generic
 - * small deterministic string/number/date helpers
 - * environment-agnostic parsing helpers
 - MUST NOT contain domain policy or application lifecycle logic.
- **api**
 - Owns all **network I/O** (HTTP/SSE/WS wrappers, uploads/downloads).
 - Normalizes errors and enforces response contracts.
- **core**

- Owns shared **non-visual logic and policies**:
 - * auth/session helpers (as policy, not routing)
 - * permission evaluation helpers
 - * shared deterministic domain utilities (prefer **foundation** for generic helpers)
 - * optional global state stores (when truly cross-feature)
 - **ui**
 - Owns reusable **presentation components** and design primitives.
 - **features/***
 - Own page-level user flows: routes, screens, feature stores, feature-local components.
 - **apps/***
 - Own composition: routing composition, bootstrap, global providers, global CSS and assets, configuration injection.
-

6 Foundational principles

6.1 P1 — Contracts first, implementation second

If multiple packages depend on a behavior, define the contract first:

- encode it as types/schemas/constants (in **contracts**)
- document semantics (edge cases, invariants)
- add at least one boundary test (contract test or stable mock)

Do not start feature implementation until the boundary contract is explicit.

6.2 P2 — Single owner per cross-cutting concern

A cross-cutting concern MUST have a single owner package. Examples:

- response envelope decoding and error mapping -> **api**
- auth/session policy helpers -> **core** (routing remains in app shell)
- permission evaluation utilities -> **core**
- reusable components and theme primitives -> **ui**
- shared DTO types and identifiers -> **contracts**

Duplicate implementations are treated as defects unless explicitly justified.

6.3 P3 — Features own user flows, not infrastructure

Feature packages MUST NOT:

- create their own HTTP clients
- implement their own global auth redirect logic
- reinterpret response envelopes

Features SHOULD:

- call `api`
- use `core/foundation` for shared logic
- use `ui` for shared components

6.4 P4 — Acyclic dependency graph, enforced

Allowed imports:

- `apps/*` -> may import from any internal package.
- `features/*` -> may import from:
 - `contracts`
 - `foundation` (if present)
 - `api`
 - `core`
 - `ui`
 - itself
- `ui` -> may import from:
 - `contracts`
 - `foundation` (preferred)
 - external UI libs
 - it MUST NOT import from `core` or from any feature.
- `core` -> may import from:
 - `contracts`
 - `foundation`
 - (rarely) `api` if truly unavoidable; prefer inversion (interfaces/events) instead.
 - NOTE: If `core` needs to reason about API failures, it MUST do so via shared error representations defined below both (`contracts/foundation`), not by importing `api`-owned error classes.
 - it MUST NOT import from `ui` or features.
- `api` -> may import from:

- `contracts`
- `foundation`
- it MUST NOT import from `core`, `ui`, or `features`.
- `contracts` -> MUST NOT import from other internal packages.
- `foundation` (if present) -> MUST NOT import from `api/core/ui/features`; it MAY import from `contracts`.

MUST: No feature-to-feature imports.

If two features need shared code, promote it into `ui`, `core`, `foundation`, or `contracts` as appropriate.

Strict rule: keep pure logic separated from external side effects To avoid “shared layer entropy” and test flakiness, SDMM distinguishes **pure logic** from **side effects**.

- **Pure logic** = deterministic computation. Given the same inputs, it returns the same outputs.
 - no network
 - no persistence
 - no global mutable state
 - no direct clock/time reads if determinism matters (inject time via `now()`)
 - no DOM access
- **Side effects** = interacting with the environment or introducing nondeterminism.

Side effects are not all equal. SDMM uses two buckets that matter for package boundaries:

1) External I/O (system effects)

- network calls (HTTP/SSE/WS)
- persistence (cookies/localStorage/IndexedDB/cache storage)
- filesystem access (Node/server runtimes)
- service worker registration
- analytics emission / telemetry exporting
- global process state and cross-tab/process coordination

2) View I/O (view effects)

- DOM APIs (measurement, focus, scroll, observers)
- browser events and event listeners
- timers and scheduling (`setTimeout`, `requestAnimationFrame`, debouncing)

Rules:

- **MUST:** Pure logic MAY be promoted into more shared, lower-level packages (`foundation` / `core`) when it stays pure.
- **MUST NOT:** External I/O MUST NOT be promoted into `foundation` / `core` / `ui`.
 - External I/O belongs in `api` (or explicitly named infra packages), orchestrated by features/app shell.
- **MAY:** View I/O is allowed in `ui` and feature packages because it is part of presentation behavior.
 - View I/O MUST be localized (component/hook scope), lifecycle-correct (cleanup on unmount), and MUST NOT run at import time.
 - View I/O MUST NOT be used as a backdoor to perform External I/O (e.g., `ui` starting background polling or writing persistent storage).

If code mixes pure logic with effects:

- split it: keep the pure core in `foundation/core`, keep the effectful adapter in `ui/feature/app shell`.

6.5 P5 — Public API only - no deep imports

Across package boundaries, imports MUST be from the package root only:

OK: `import { fetchCourses } from '@org/api'`

NOT OK: `import { fetchCourses } from '@org/api/src/course/fetch'`

Define the public surface explicitly Each package MUST have a single entry (e.g., `src/index.ts`) exporting:

- stable public functions/types/components
- nothing internal

If you intentionally expose multiple public entrypoints (subpaths), they MUST be declared as public API (e.g., via export maps) and treated as contracts.

6.6 P6 — Backward-compatible behavior is the default

During modularization, runtime behavior SHOULD remain stable by default.

Behavior changes MUST be treated as product changes:

- documented
- reviewed
- tested

6.7 P7 — Make packages small, cohesive, and replaceable

A package SHOULD be:

- owned by a clear team/person
- understandable with local context
- replaceable without rewriting the whole repo

If a package grows beyond “AI-safe” size, split it by responsibility, not by file count.

6.8 P8 — The app shell is the composition root

MUST: Only the app shell wires together:

- router + route guards
- global providers (auth context, query client, i18n)
- infrastructure configuration (API base URLs, token providers, loggers)
- global error handling and notifications policy
- telemetry setup

MUST NOT: Packages create app-global behaviors at import time (e.g., registering router guards, installing plugins, starting timers, mutating global singletons), unless explicitly coordinated through the app shell.

MUST: If a package needs global registration, it MUST expose an explicit `registerX(appContext)` (or `install(app)`) function, and the app shell MUST call it during bootstrap.

6.9 P9 — Prefer explicit dependency injection over hidden globals

SHOULD: Prefer explicit construction over implicit module singletons.

Recommended approach (best balance of ergonomics and testability):

1) Implement a factory:

- `createX(config) -> X` (pure construction; no hidden globals)

2) (Optional convenience) Provide a default instance configured by the app shell:

- `configureX(config)` sets a default instance
- exports that rely on the default MUST fail fast if not configured

MUST: If you provide a `configureX()` convenience singleton, you MUST also provide a factory (`createX()`).

MUST: Libraries/packages SHOULD NOT require global configuration to be import-safe.

MUST: Tests MUST be able to create isolated instances without relying on global mutable module state.

If you choose a configured default instance pattern:

- configuration MUST be single-assignment in production (fail fast on reconfigure)
 - a `resetForTests()` hook MAY exist, but it is not sufficient if tests run concurrently in the same process
 - in test-intensive or concurrency-heavy environments, prefer instance-based injection (factory + explicit passing) as the default
-

7 Contract rules

7.1 Contract sources of truth

MUST: For each contract, identify exactly one source of truth:

- backend-owned API contracts -> backend schema/spec (mirrored into `contracts`)
- frontend-owned identifiers (route IDs, permission codes) -> `contracts` + app shell routing
- UI component contracts -> ui public API (props/events) + story/tests

7.2 Contract evolution: refactors vs versioning

The word “versioning” is overloaded; SDMM uses two modes:

1) Refactor-mode (single deploy, single change set)

When all consumers are updated together (typical within one repo + one app build), treat contract changes as **refactors**:

- prefer additive changes
- update all usages in the same change set
- keep behavior stable unless intentionally changed

2) Versioning-mode (independent deploy cadence or external consumers)

When consumers cannot be updated atomically (separate services, multiple deploy pipelines, third-party clients), treat contract changes as **versioned**:

- introduce explicit v2/v3 (URL versioning or header/content negotiation)
- maintain a compatibility window
- measure adoption and deprecate intentionally

MUST: A boundary that crosses independent deploy lifecycles MUST use versioning-mode.

7.3 Contract drift prevention

MUST: Prevent silent drift by enforcing at least one of:

- contract tests at boundaries

- runtime schema validation at boundaries for untrusted inputs (when practical)
 - shared mocks generated from `contracts` and used by features
-

8 api package guidelines (I/O layer)

8.1 Responsibilities

`api` owns:

- HTTP client configuration (base URL, interceptors, headers)
- request/response decoding (including envelopes)
- typed endpoint wrappers grouped by domain
- error normalization (mapping raw transport/library errors into shared error representations defined in `contracts` and/or `foundation`)
- upload/download helpers
- streaming wrappers (SSE/WS) where applicable
- request dedup/caching policy if used

8.2 Configuration injection

Because `api` SHOULD NOT import app-specific state (router/stores), it MUST be configurable from the app shell.

Recommended pattern:

- `api.configure({ baseUrl, getAuthToken, onAuthError, onError, logger, now })`

Where:

- `getAuthToken` returns a token (or `undefined`) and MAY be async
- `onAuthError` handles auth failure at the app level (redirect/reset)
- `onError` is an optional hook for centralized error reporting
- `now` is an injectable clock for testability

MUST: `api` must be usable in tests without a browser/router/store by replacing hooks with test doubles.

Reference implementation shape (illustrative):

```
export type Awaitable<T> = T | Promise<T>

export type GetAuthToken = () => Awaitable<string | undefined>

export interface ApiConfig {
```

```

baseUrl: string
getAuthToken?: GetAuthToken
onAuthError?: (err: unknown) => void
onError?: (err: unknown) => void
now?: () => number
}

export interface ApiClient {
  request<T>(path: string, init?: RequestInit): Promise<T>
  // Add domain wrappers on top, or expose typed per-domain modules.
}

// 1) Factory (preferred for tests and isolation)
export function createApiClient(cfg: ApiConfig): ApiClient {
  // return an instance that closes over cfg; no hidden globals
  throw new Error('implement')
}

// 2) Convenience singleton (for app shell ergonomics)
let _defaultClient: ApiClient | undefined
let _configured = false

export function configureApi(cfg: ApiConfig): void {
  if (_configured) {
    throw new Error('configureApi called more than once')
  }
  _defaultClient = createApiClient(cfg)
  _configured = true
}

export function api(): ApiClient {
  if (! _defaultClient) {
    throw new Error('API not configured. Call configureApi() in app shell bootstrap.')
  }
  return _defaultClient
}

// Optional: only for unit tests that must use the singleton.
export function resetApiClientForTests(): void {
  _defaultClient = undefined
}

```

```
    _configured = false  
}
```

MUST: If `getAuthToken` can refresh tokens, it MUST implement single-flight behavior (avoid stampedes) and it MUST be safe to call on every request.

MUST: Configure `api` during application bootstrap (in the app shell) before any feature invokes endpoints.

8.3 Timeouts, cancellation, and long-running operations

MUST: Every request has an explicit timeout policy (global default + per-endpoint overrides).

Classify operations:

- **Interactive** (user waiting): strict time budget, clear loading UX, cancellation when possible
- **Background** (user can navigate away): job-based patterns, resumability
- **Compute/generation/media**: explicit long timeout, progress hints, idempotency keys, best-effort cancellation

MUST: If the UI can trigger the same long-running action repeatedly, define dedup/idempotency strategy explicitly.

8.4 Error model ownership

A common SDMM failure mode is placing “shared” error classes in `api`, then writing retry/auth/session policies in `core` that need to inspect those classes.

That creates one of two bad outcomes:

- `core` imports `api` (violating the intended purity boundary), or
- developers duplicate error representations across packages (drift and inconsistent handling).

MUST: Define the *shape* of normalized errors **below both** producers and consumers:

- Prefer contracts for the **data representation** (types, enums, schemas).
- Use `foundation` (optional) for **pure runtime helpers** (constructors, type guards) if needed.

`api` owns the **mapping logic**:

- transport/library error -> normalized error representation
- API envelope failure -> normalized error representation

`core` owns **pure policies** over the normalized representation:

- retry classification: `shouldRetry(err)`
- session/auth classification: `isSessionExpired(err)`
- user messaging policy inputs (but not UI rendering)

MUST NOT: Require `core` (or `ui`) to import error classes/types that are defined in `api`. If an error representation must be shared, move it down into `contracts/foundation`.

SHOULD: Prefer discriminated unions (or tagged objects) over `instanceof` across package boundaries.

Example (illustrative):

```
// contracts/errors.ts (representation only)
export type NormalizedError =
  | { kind: 'auth'; status?: number; code?: string; message: string }
  | { kind: 'transport'; message: string; retryAfterMs?: number }
  | { kind: 'api'; code: string; message: string }
  | { kind: 'cancelled'; message?: string }
  | { kind: 'unknown'; message: string; detail?: unknown }

// core/policy.ts (pure)
export function shouldRetry(e: NormalizedError): boolean {
  return e.kind === 'transport'
}

export function isSessionExpired(e: NormalizedError): boolean {
  return e.kind === 'auth' && (e.status === 401 || e.status === 403)
}

// api/errors.ts (I/O layer: mapping only)
export function normalizeError(raw: unknown): NormalizedError {
  // map axios/fetch/etc to NormalizedError
  return { kind: 'unknown', message: 'Unexpected error', detail: raw }
}
```

This pattern keeps the dependency graph acyclic and allows policy evolution without leaking I/O concerns upward.

9 contracts package guidelines (contract layer)

9.1 What belongs here

`contracts` SHOULD contain:

- DTO types: request/response shapes, shared domain models
- error code enums and typed error payload shapes
- normalized error representations used across packages (as pure data types), when you need shared retry/auth policies
- identifiers: route IDs, permission codes, feature flags
- event/message schemas for streaming
- optional runtime schemas for boundary validation

MUST: Keep `contracts` as representations only.

Behavior like retry logic, auth redirect behavior, caching policy, or envelope parsing MUST NOT live here.

If you need to standardize semantics, encode them as:

- documented rules (this document)
- tested behavior in `api/core`
- constants that parameterize behavior (only if truly stable)

9.2 What must not belong here

`contracts` MUST NOT contain:

- UI components
- network clients
- environment-specific logic (reading `window/document` or process env directly)
- orchestration code

Goal: `contracts` is boring, stable, and dependency-light.

10 core package guidelines (shared policy + pure logic layer)

10.1 Responsibilities

`core` owns:

- auth/session policy helpers (not routing)
- permission evaluation helpers
- cross-feature non-visual logic and policies
- shared stores only when truly global

10.2 Keep `core` deterministic by default

SHOULD: `core` should be I/O-free by default.

If `core` must perform External I/O, isolate it behind explicit interfaces and keep it rare.

10.3 Framework specificity

`core` SHOULD avoid coupling to specific UI frameworks.

If framework hooks are needed, isolate them into a thin adapter submodule and keep the majority framework-agnostic.

11 ui package guidelines (shared presentation layer)

11.1 Responsibilities

`ui` owns:

- reusable presentation components
- design tokens/theme primitives (as code and/or CSS variables)
- shared layout blocks (when they are genuinely reusable)
- component-level accessibility patterns

11.2 Rules (avoid hidden coupling)

`ui` components SHOULD be:

- prop-driven
- slot/children-friendly
- stable in public API
- free of feature-specific state

`ui` MUST NOT import feature stores or feature logic.

11.3 Dependency rule: do not couple ui to core

SHOULD: Treat `ui` as “dumb” presentation.

- `ui` MUST NOT depend on `core` as a default rule.
- If `ui` needs generic helpers (formatting, parsing), prefer `foundation`.
- If a `ui` component needs policy decisions (permissions, routing, auth state), pass decision-s/results as props from the feature/app shell.

If you insist on `ui -> core` in a specific repository:

- you MUST prove `core` is I/O-free and policy-light in the imported area
- you MUST enforce boundaries to prevent `core` from importing `ui` (cycle risk)
- you SHOULD treat this as an exception, not a norm

11.4 Styling and assets strategy (required)

Modular frontends fail most often at styling boundaries. Define this explicitly.

MUST: Choose one primary scoping strategy and enforce it:

- **Scoped styles / CSS Modules:** preferred for package isolation
 - feature-local and UI component styles MUST be scoped
 - global selectors are disallowed except in app shell global stylesheet
- **Utility-first CSS (e.g., Tailwind):** acceptable if centrally configured and shareable across packages
 - the canonical config/preset MUST live in a shared place (e.g., `ui` or a dedicated `design-tokens` package, or repo root), not only inside an app shell
 - app shells and package tooling (storybook/tests) MUST consume the same preset/-config
 - avoid per-package divergent configs; allow only additive extension in app shell when necessary
- **CSS-in-JS:** acceptable if SSR constraints and runtime costs are understood

MUST: Define ownership:

- global resets, typography baseline, and “page chrome” styles -> app shell
- design tokens (colors, spacing, typography scale) -> `ui` (or `foundation` if purely non-visual constants)
- component styles -> `ui` or feature-local components, but always scoped

MUST: Prevent style collisions:

- no package is allowed to inject global CSS at import time (except app shell)
- class naming conventions or scoping MUST be enforced (lint/build rules)

MUST: Define static asset ownership:

- shared icons/fonts/images used across features -> `ui` (or app shell if truly global)
 - feature-specific assets -> feature package
 - assets MUST be imported via stable public APIs (no deep imports across packages)
-

12 Feature package guidelines (`features/*`)

12.1 Feature boundaries

Each feature package MUST be “complete” for its domain:

- it can be removed and the app still builds (routes removed)
- it does not require another feature to compile

12.2 Feature public exports

Feature packages SHOULD export only:

- routes/screens registration
- optional menu descriptors (if your app composes menus from features)
- minimal feature-level constants/types

If feature internals are reused by other features, they likely belong in `ui`, `core`, `foundation`, or `contracts`.

13 App shell guidelines (`apps/*`)

The app shell owns:

- route composition and guards
- provider composition
- global CSS and assets
- configuration injection for `api` and other infra
- top-level error boundaries and telemetry wiring

Feature packages MUST NOT register global guards or global plugins.

14 State management principles

14.1 Where state lives

- truly global state -> `core` or app shell
- feature state -> feature package

14.2 Avoid hidden coupling

Features MUST NOT read/write other feature stores directly.

Shared state must be promoted into `core` (or persisted to the backend).

15 Performance and build efficiency principles

15.1 Guarantee singleton framework/runtime dependencies

MUST: Prevent duplicate runtime copies:

- align versions via workspace constraints
- prefer peer dependencies for framework libs in packages that should not bundle them
- use lockfile enforcement/resolutions to avoid drift
- run a bundle/dependency duplication check regularly (CI preferred)

15.2 Deliberate code splitting

SHOULD: Code-split by routes/screens:

- feature pages SHOULD be lazy-loaded
- shared packages SHOULD avoid pulling heavy dependencies into the initial chunk

15.3 Export hygiene (avoid barrel bloat)

MUST: Do not export “everything” from root barrels by default.

Exports are part of your contract; keep them explicit and minimal.

15.4 Build/CI efficiency

SHOULD: Use caching and incremental tasks:

- per-package typecheck/tests
- build caching (Nx/Turbo or CI cache) where available

16 Error handling, security, and observability

16.1 One global error policy

MUST: Centralize user-visible error handling policy.

Features may catch and recover, but must not invent new inconsistent policies.

16.2 Correlation IDs and telemetry

SHOULD: Propagate correlation IDs through:

- client logs
- API requests (headers)
- error reports

16.3 Secret handling

MUST: No secrets in repo source code.

Use environment variables and secret managers.

16.4 Untrusted inputs and boundary validation

SHOULD: Validate untrusted inputs at boundaries (API responses, query params, storage).

Use runtime schemas where the risk justifies the cost.

17 Testing strategy

17.1 What to test where

- **contracts:** schema/type-level tests and example fixtures
- **api:** client behavior tests (envelopes, error mapping, retry/timeout policy)
- **core:** policy logic tests (permission evaluation, session helpers)
- **ui:** component behavior tests (props/slots/accessibility)
- **features:** integration-ish tests of user flows
- app shell: smoke tests + routing/guard behavior

17.2 Contract tests

SHOULD: Add contract tests where failures are expensive:

- response decoding and error semantics
 - streaming lifecycle correctness
 - permission/route identifier stability
-

18 Tooling enforcement

Minimum enforcement set:

- path aliases for package root imports
 - lint rules for boundary imports (no deep imports, no feature-to-feature)
 - typecheck gate at workspace root
 - build gate for each app shell
 - (recommended) automated duplicate dependency/bundle check
-

19 AI-assisted development workflow

When delegating to AI, every task MUST include:

- 1) **Scope**: which packages may be modified
- 2) **Contracts**: which identifiers/shapes/semantics must remain stable
- 3) **Acceptance criteria**: build/typecheck/test expectations
- 4) **Non-goals**: what must not be refactored

AI agents perform best when boundaries, public APIs, and contracts are explicit.

20 Decision guide

20.1 S1 — Add a new screen/page

- implement in a feature package
- export route registration
- app shell composes routes

20.2 S2 — Add a new backend/API endpoint

- add typed wrapper in `api`
- export from `api` root
- features consume wrapper only

20.3 S3 — Add a shared UI component

- used in one feature -> keep local
- used across features -> promote to `ui`

20.4 S4 — Add shared logic/utility

- generic/pure -> `foundation`
- cross-feature policy -> `core`

20.5 S5 — Add a new permission or gated UI action

- define stable identifiers in `contracts`
- evaluate in `core` or app shell
- render gating in features/`ui` via props or directives/hooks

20.6 S6 — Add a streaming/realtme feature

- define event schemas in `contracts`
- implement lifecycle wrapper in `api`
- features handle presentation only

20.7 S7 — Add a long-running compute/generation operation

- implement request in `api` with explicit timeout/cancellation
- use idempotency keys if supported
- UI shows progress and prevents accidental duplicates

20.8 S8 — Two features need the same domain type

- if it crosses package boundaries, put it in `contracts`

20.9 S9 — Feature A needs data produced by Feature B

Prefer:

- 1) backend as source of truth
- 2) promote shared state to `core`

Avoid direct feature-to-feature imports.

20.10 S10 — Change a contract (shape/semantics/identifier)

Treat as high-risk:

- update `contracts`
- update implementations (`api/core`)
- update consumers
- add/adjust boundary tests

20.11 S11 — A feature package becomes too large

Split by responsibility (sub-features) or promote reusable parts into `ui/core/foundation`.

20.12 S12 — You need to break a dependency rule

Only via an explicit exception process:

- document why
- define sunset plan
- add automated check to prevent spread

Appendix A. Migration playbook

- 1) Extract **contracts** (types/identifiers) without behavior change
 - 2) Extract **api** (clients + wrappers) without behavior change
 - 3) Extract **ui** shared components
 - 4) Extract **core** policies and truly global stores
 - 5) Split features one-by-one (routes + pages + feature stores)
 - 6) Add enforcement tooling gates (lint/typecheck/build)
 - 7) Add boundary/contract tests for critical paths
 - 8) Add bundle duplication checks and code-splitting validation
-

21 SDMM checklist

- The product still builds and deploys as a **single application artifact** per app shell.
- No duplicate framework/runtime copies in the output.
- CI includes an automated duplicate dependency/bundle check (or a periodic enforced check) to catch drift early.
- Shared design-system config (tokens, Tailwind preset, etc.) is owned by **ui/design-tokens** and consumable by app shells and UI tooling without **ui -> app** imports.
- All network I/O is centralized in **api** (or explicitly named infra packages).
- No feature imports another feature.
- Shared UI lives in **ui**; shared pure helpers live in **foundation**; shared policies live in **core**; shared contracts live in **contracts**.
- Contracts are documented and encoded in types/schemas (validated where needed).
- Streaming connections and long-running operations have explicit lifecycle handling.
- Automated checks enforce boundary rules (lint/typecheck/build gates).

⁰© 2026 AzzCraft, Inc. (艾之舟科技) All rights reserved. This document is published as a public reference. Commercial use, redistribution, or derivative works require prior written permission. For licensing inquiries, contact qipeng@azzcraft.com.